

Notebook to Cloud



Yue Sun



Dylan Randle



Bhaven Patel



DH Lee

Workshop Overview for Session 1

Intro Lecture
9:45 - 10:15
David



Code Performance
10:30 - 11:10
Yue



Containers
11:10 - 12:10
Dylan



Hands-on Kubernetes
1:30 - 2:30
Bhaven & DH



Lunch
12:30 - 1:30



Set up Kubernetes
12:10 - 12:30
Bhaven

Why Notebooks?

- Easy to use
- Great for prototyping
- Excellent for documentation and examples

We want to find where the functions $y_1 = x$ and $y_2 = \exp(-2\sin^2(4x))$ intersect. That is, for what values of x is $y_1 = y_2$? This question is equivalent to the problem of finding the zeros of

$$f(x) = x - \exp(-2\sin^2(4x)).$$

Before doing anything, it's a good idea to visualize what we're up against.


In [2]: Slide Type Sub-Slide ▾

```
# Define function
x = np.linspace(0.0, 2.0*np.pi, 1000)
y = np.exp(-2.0 * np.sin(4.0*x)*np.sin(4.0*x))
```


In [3]: Slide Type Sub-Slide ▾

```
# Plot the two functions to see where they intersect
fig, ax = plt.subplots(1,1, figsize=(10,6))
ax.plot(x, y, lw=3, label='y1 = \exp\left(-2\sin^2\right)\left(4x\right)\right')
ax.plot(x, x, ls='--', lw=3, label='y2 = x')

ax.set_xlim(0, 2.0)
ax.set_ylim(0, 2.0)
ax.set_xlabel('x', fontsize=24)
ax.set_ylabel('y', fontsize=24)
ax.set_title('Intersections of y1 and y2', fontsize=24)
ax.tick_params(labelsize=24)
ax.legend(fontsize=24)
plt.tight_layout()
```

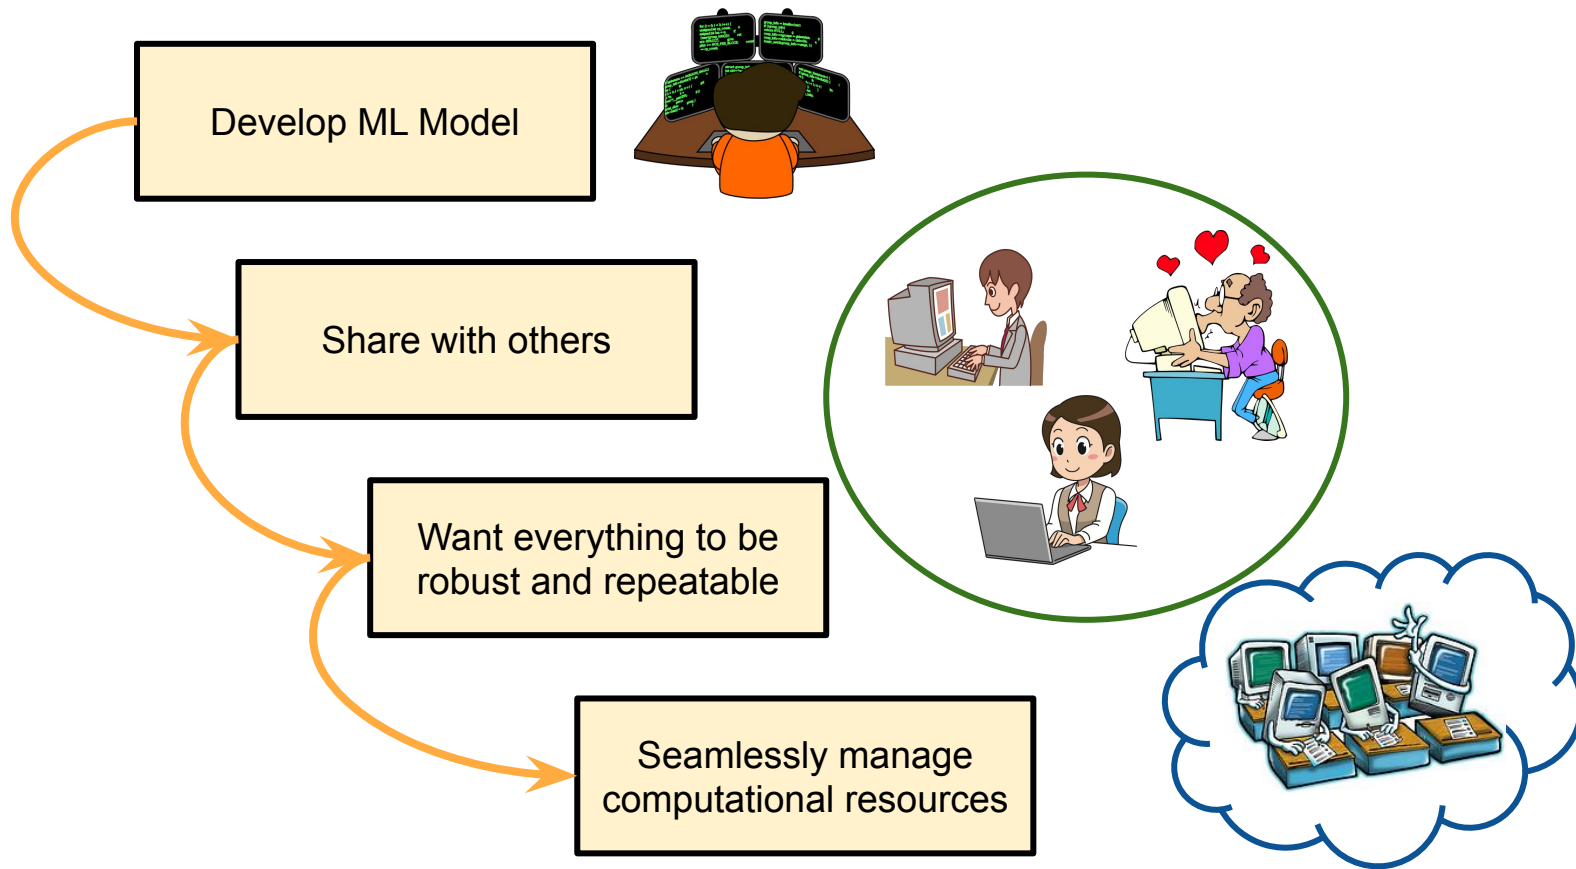


Why Cloud?




- *Many* resources beyond what you have locally
 - CPUs
 - GPUs
 - Storage
 - ...
- Resources maintained by experts
- Easier to reproduce results
 - Collaborators have replicable platforms
- Easier to host applications
 - Virtual machines
 - Containers
 - ...

A Dream Scenario




A Common Scenario


Develop code base in
Jupyter notebooks



Pass around to
colleagues and friends



Tweak depending on versions and
platforms people are using




A Modern Approach




Prototype, document,
and examples in
Jupyter notebooks

1




Large code-bases
written in scripts




Managed with
version control

2



Ship package in
lightweight containers

3



Deploy to
the cloud

4



Automatically manage
containers in cloud

5

Towards Collaboration: What Can Go Wrong

- Congrats! You've developed a nice application.
- But you can't work in isolation:
 - Need help and input from colleagues
 - Want to have other people use your awesome app
- What can go wrong?
- Turns out, a lot!
 - Different software versions (e.g. Python 2.7 vs. 3.7)
 - Different operating systems (Windows vs. Linux vs. MacOS)
 - Different packages required than are available locally

Achieving Collaborative Isolation

- We develop on specific platforms using specific versions of software and dependencies
- How can we make sure everyone works with the same environment?
 - Virtual environment --- Still depends on operating system
 - Virtual machines --- Full isolation
 - Containers --- Only virtualize OS (not the hardware)

Virtual Environments



- ✓ Requirement A
- ✓ Requirement B
- ✓ Requirement C
- List requirements in special file
- Automatically install dependencies




- Create virtual environment
- Install package
- Dependencies installed too
- ✓ Requirement A
- ✓ Requirement B
- ✓ Requirement C

Comments on Virtual Environments


- Many options for virtual environments
 - virtualenv
 - conda
- Very useful for working with fellow developers
- Convenient, lightweight method to achieve code portability
- Not as helpful when deploying a package to a larger audience
- They do not provide complete isolation
 - Still depends on your operating system
 - Uses global packages and dependencies of the operating system

The Other End of the Spectrum: Virtual Machines



- A virtual machine (VM) is a file that acts like a separate computer system.
- You can install a completely different operating system on this virtual machine.
 - e.g. You run MacOS; create a VM and install Windows on it
- VMs have their own virtual hardware
 - CPUs, memory, hard drives, etc.
- Software inside the VM can't affect the actual computer
 - Sandboxing
 - Safe place to test virus-infected software

Comments on Virtual Machines

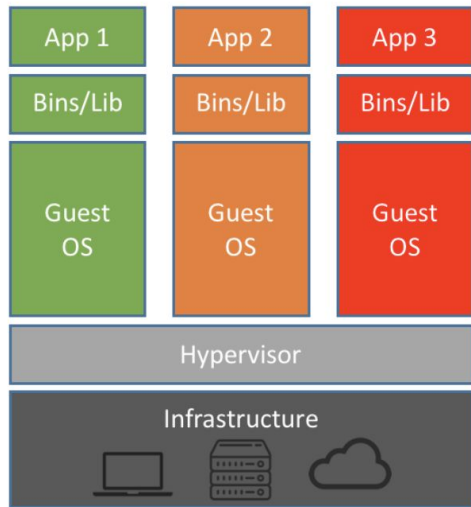
- They can help lower costs and can be more efficient
 - e.g. no need to spend money on physical hardware and cooling systems
 - Best Virtual Machines of 2019
 - VMware
 - VirtualBox
 - There is overhead associated with VMs
 - They may not be as fast as the host system
 - They may not have the same graphics capabilities
 - They can take some time to start up (order of minutes)
- 
- Powerful but heavyweight

Amazon Web Services
(AWS)

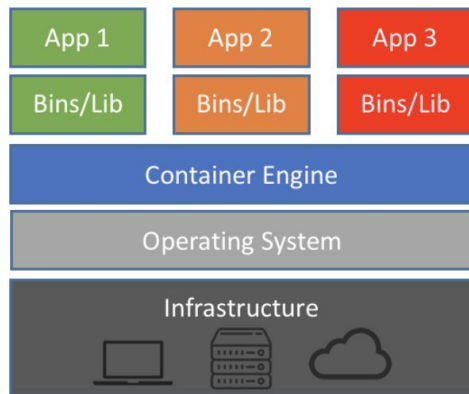
Google Cloud

Microsoft Azure

Containers



Machine Virtualization



Containers

- Only virtualize the OS
- Give impression of separate OS
 - Much cheaper than VMs
- e.g. Create container on Mac, but install Linux OS
 - Container still works on Mac
 - Inside container it's like Linux
- Benefits:
 - Lightweight
 - Quick start-up time
 - Pseudo-isolation
 - Run many at once on a system

To the Cloud: Managing Containers with Kubernetes



- What if you have a bunch of containers?
 - Either working together or
 - Independent but taking up resources or
 - Both
- K8s does all of the container management for you!
- K8s is an open-source platform for container management developed by Google.
- It allows users to define rules for how container management should occur, and then it handles the rest!

There is so much more...

- We're leaving a lot on the table today
 - Containers for high performance computing: Singularity
 - Amazon Sagemaker : Fully automated machine learning service
 - Much, much more
- Goals for today:
 - Perspective on efficient development practices
 - Best practices for deploying models - containerization
 - Deployment to the cloud and container management
 - First contact with AWS
- DH will summarize the big ideas at the very end

Workshop Overview for Session 1

Intro Lecture
9:45 - 10:15
David



Code Performance
10:30 - 11:10
Yue



Containers
11:10 - 12:10
Dylan



Hands-on Kubernetes
1:30 - 2:30
Bhaven & DH



Lunch
12:30 - 1:30



Set up Kubernetes
12:10 - 12:30
Bhaven